EnergyMind – AI Agents for Smart Energy Insights

Course Name - Agentic AI

Institution Name - Medicaps University – Datagami Skill Based Course

:

Sr.No.	Student Name	Enrollment Number
1.	Srashti Bhadoriya	EN22CS301982
2.	Pranjal Sharma	EN22CS301731
3.	Riya Patidar	EN22CS301813
4.	Ritik Patidar	EN22CS301809
5.	Roshni Patidar	EN22CS301825

Group Name - Group 10D7

Project Number - AAI - 33

Industry Mentor Name - Aashruti Shah

University Mentor Name - Ajeet Singh Rajput

Academic Year -2025–26

Problem Statement

The Energy sector plays a critical role in global economic stability and sustainable development. Rapid advancements in renewable energy technologies, fluctuations in oil and gas markets, climate policy reforms, and energy storage innovations require continuous and structured research. However, manual research processes are time-intensive, fragmented across multiple sources, and often lack consistency in analysis and presentation.

Conventional AI chatbots provide general responses but do not follow a structured workflow. They lack domain validation, iterative quality control, and persistent session memory. As a result, outputs may be unstructured, lack depth, or include irrelevant information outside the intended domain.

To address these challenges, this project proposes the development of an **Autonomous Energy Researcher Agent**, a multi-agent AI system capable of performing domain-restricted, real-time research, structured report generation, internal quality review, and secure data management within the energy industry.

Manual research in this domain involves:

- Searching multiple sources.
- Validating authenticity.
- Extracting relevant insights.
- Structuring findings into professional reports.

Traditional AI chatbots provide one-time responses without structured workflows, domain validation, or quality control. They:

- Do not verify if queries belong to a specific industry.
- Lack iterative refinement mechanisms.
- Cannot maintain structured session memory.
- Often produce unverified or unstructured outputs.

Therefore, this project aims to design and implement an **Autonomous Energy Research Agent** that performs real-time research, applies structured analysis, validates quality through a review loop, and maintains session-level memory while ensuring domain restriction and user authentication.

Project Objectives

The primary objective of this project is to design and implement a domain-specific autonomous AI system that automates research tasks within the energy sector. The system leverages LangGraph to orchestrate multiple AI nodes in a structured workflow and uses FastAPI to build secure backend APIs. The frontend interface is developed using Streamlit to enable user interaction.

The main objectives of the project are:

- To design a structured multi-agent workflow using LangGraph.
- To integrate real-time web search using Tavily AI for up-to-date information.
- To implement domain validation using a Gatekeeper node.
- To generate structured research reports including executive-level insights.
- To implement a Reviewer node that evaluates report quality and triggers revisions.
- To maintain short-term conversational memory using thread-based state management.
- To develop secure REST APIs using FastAPI.
- To build an interactive frontend using Streamlit.
- To implement user authentication using JWT and bcrypt.
- To implement caching and database storage for performance optimization.
- To deploy the complete system using Docker and cloud hosting.

Scope of the Project

The scope of this project is limited to automated research within the Energy Industry. The system supports analysis of renewable energy trends, oil and gas market developments, energy storage technologies, climate policies, and sustainability initiatives. It generates structured reports containing executive summaries, market trends, opportunities, risks, and future outlook.

The project includes secure user authentication, chat history storage, database-based caching, and file-based report persistence. However, the system does not currently support multi-domain research, enterprise-scale distributed databases, or long-term semantic vector memory.

The scope of this project covers:

- Real-time automated research in the Energy Industry.
- Multi-stage AI-based analysis pipeline.

- Structured professional report generation.
- Query caching using a Knowledge Base database.
- User-based chat history management.
- Secure login and signup functionality.
- Cloud-based deployment for public accessibility.

The system is specifically restricted to the Energy domain, including:

- Renewable energy.
- Oil and gas.
- Energy storage technologies.
- Climate policy.
- Electricity markets.
- Sustainability initiatives.

PROPOSED SOLUTION

The proposed solution, **EnergyMind AI**, is a full-stack autonomous research system built on a multi-agent architecture. Rather than producing a single direct response like a conventional LLM application, the system simulates a structured team of AI agents that collaborate to complete complex research tasks. Each agent is responsible for a specific function in the workflow, enabling modularity, scalability, and improved reliability.

The backend of the system is implemented using **FastAPI**, ensuring high performance and efficient API handling. The frontend interface is developed with **Streamlit**, providing an interactive and user-friendly dashboard for submitting research queries and viewing structured reports. Workflow orchestration is managed through LangGraph, which enables controlled multi-step execution and state management between agents. Large language model inference is powered by Groq, ensuring fast and optimized response generation.

The system processes user queries through a structured pipeline that includes validation, research, analysis, report generation, and review. This layered architecture enhances accuracy, maintains contextual relevance, and ensures that the final output is well-structured and logically organized.

Key Components of the Proposed Solution

1. Multi-Agent Architecture

- Simulates a collaborative AI team
- Each agent performs a specialized task
- Improves modularity and maintainability

2. Backend (FastAPI)

- Handles API requests and routing
- Ensures high performance and scalability

3. Frontend (Streamlit)

- Provides interactive UI
- Displays structured research reports

4.Workflow Orchestration

- Managed using LangGraph
- Controls agent execution flow
- Maintains state across research stages

5. LLM Inference

- Powered by Groq
- Enables fast and efficient response generation

6. Processing Stages

- Query validation
- Information research
- Data analysis
- Structured report generation
- Final review and refinement

This design ensures that EnergyMind AI produces accurate, comprehensive, and well-structured research outputs through a systematic and autonomous workflow.

Key Features

The EnergyMind AI system incorporates a set of well-defined core features that ensure reliability, scalability, and high-quality research output. Each feature is designed to support the autonomous, multi-agent workflow and enhance the overall efficiency of the system.

1. Domain Validation through Gatekeeper Node

- A dedicated Gatekeeper agent validates whether the user query falls within the defined energy domain.
- Prevents irrelevant or out-of-scope queries from entering the workflow.
- Reduces unnecessary computation and improves overall system precision.

2. Real-Time Information Retrieval

- Integrates live search tools to fetch the latest market trends, policies, and industry updates.
- Overcomes LLM knowledge cutoff limitations.
- Ensures reports are based on up-to-date and reliable information sources.

3. Multi-Stage AI Workflow Architecture

- Structured pipeline consisting of Planner, Researcher, Analyst, and Reviewer nodes.
- Each node performs a specialized function in a sequential and controlled flow.
- Orchestrated using LangGraph for state management and execution control.

4. Structured Executive Report Generation

- Generates professional, well-formatted research reports.
- Includes sections such as executive summary, key insights, data analysis, and conclusions.
- Designed for industry-level presentation and decision-making support.

5. Automated Quality Review and Refinement Loop

- A dedicated review agent re-evaluates the generated output.
- Ensures logical consistency, completeness, and clarity.
- Refines content iteratively before delivering the final report.

6. Session-Based Memory Management

- Maintains conversation context within a user session.
- Enables follow-up queries without losing previous state.
- Improves coherence and personalization of responses.

7. Secure User Authentication and Authorization

- Implements role-based access control mechanisms.
- Ensures that only authorized users can access research tools and reports.
- Protects sensitive query data and system resources.

8. Database Caching for Repeated Queries

- Stores previously processed queries and results in a database.
- Reduces redundant API calls and computation cost.
- Improves response time for frequently requested topics.

9. Docker-Based Deployment Support

- Containerized architecture for easy deployment across environments.
- Ensures consistency between development and production setups.
- Enhances portability, scalability, and system reliability.

These features collectively make EnergyMind AI a robust, scalable, and production-ready autonomous research platform capable of delivering accurate and structured energy sector insights.

Overall Architecture / Workflow

EnergyMind AI follows a **Client–Server architecture**, where responsibilities are clearly separated between the user interface, backend services, orchestration layer, and database. The complete flow is:

User → Streamlit Frontend → FastAPI Backend → LangGraph Workflow → Database → Response

This architecture ensures modularity, scalability, and controlled multi-stage processing of research queries.

Detailed Workflow

1.User Query Submission

- The user submits a research query through the Streamlit interface.

2.API Communication

- The frontend sends a POST request to the FastAPI backend.

3. Cache Verification

- The backend checks the KnowledgeBase database table for previously stored (cached) results.
- If a cached response exists, it is directly returned to reduce computation time.

4. Workflow Triggering

- If no cached result is found, the multi-agent workflow is triggered using LangGraph.

5. Gatekeeper Validation

- The Gatekeeper node verifies whether the query belongs to the Energy domain.
- Out-of-domain queries are rejected at this stage.

6.Live Research Phase

- The Researcher node performs real-time search using the Tavily API to gather current information.

7. Analytical Processing

- The Analyst node processes collected data to identify key trends, risks, and opportunities.

8. Report Generation

- The Writer node generates a structured research report with proper sections and insights.

9. **Quality Review Loop**

- The Reviewer node evaluates the report for structure, depth, and logical consistency:

  1. **If FAIL** → Feedback is sent back to the Writer for refinement.
     1.
  2. **If PASS** → The workflow proceeds to the next stage.

10. **Follow-up Suggestions**

- A Suggestions node generates additional related research questions for deeper exploration.

11. **Data Persistence**

The final result is:

- Saved in the KnowledgeBase database.
- Stored in ChatHistory for authenticated users.
- Saved as a text file for backup purposes.

12. **Response Delivery**

- The final structured output is returned to the frontend and displayed to the user.

Tools & Technologies Used

EnergyMind AI is developed using a modern full-stack technology stack that supports scalable backend processing, interactive frontend design, secure authentication, and AI-driven workflow orchestration. The system integrates LLM inference, real-time search capabilities, database management, and containerized deployment to ensure production-level reliability and performance.

Below is a structured overview of the technologies used and their purpose:

Core Development Stack

- **Python** – Primary programming language used for backend logic, workflow implementation, and AI integration.
- **FastAPI** – High-performance REST API framework used to build backend services and handle HTTP requests efficiently.
- **Streamlit** – Frontend framework used to create an interactive and user-friendly research interface.

AI & Workflow Components

- **LangGraph** – Workflow orchestration framework used to manage multi-agent execution flow and state transitions.
- **Groq (Llama-3.1-8B)** – LLM inference engine used for generating analytical reports and structured responses.
- **Tavily AI** – Real-time search API used to retrieve up-to-date energy sector information.

Database & Security

- **SQLAlchemy** – ORM (Object Relational Mapper) used to interact with the database using Python models instead of raw SQL queries.
- **SQLite** – Lightweight relational database used to store cached reports, chat history, and user data.

- **JWT (JSON Web Token)** – Implements secure token-based authentication for user sessions.
- **bcrypt** – Used for secure password hashing and storage.

Deployment & DevOps

- **Docker** – Containerization platform used to ensure consistent deployment across different environments.
- **Hugging Face Spaces** – Cloud hosting platform used for deploying and demonstrating the application online.

RESULTS & OUTPUT

- The **EnerGpt** home screen serves as the primary entry point for energy research. It features a minimalist dark-themed UI with a central search bar for natural language queries and a sidebar for navigation, session management, and guest access.

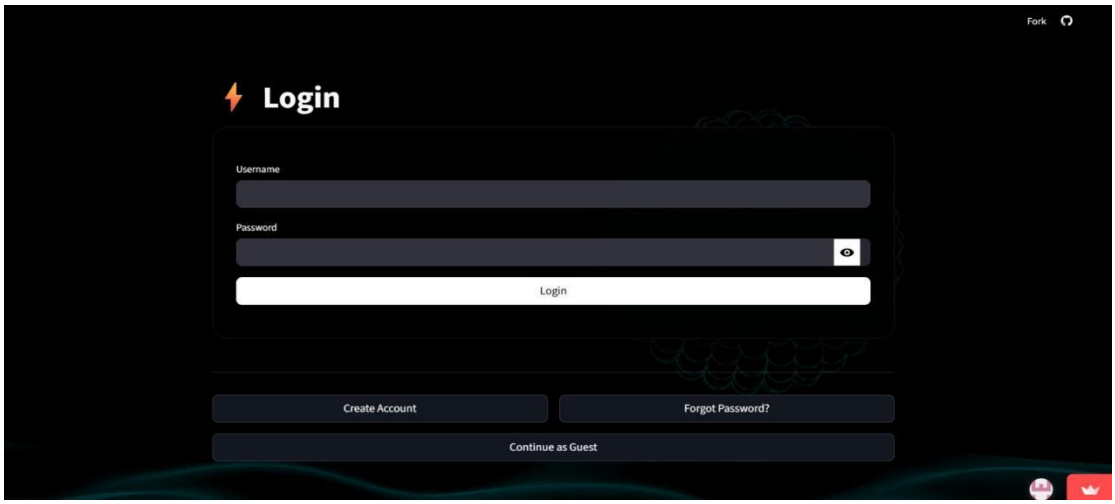

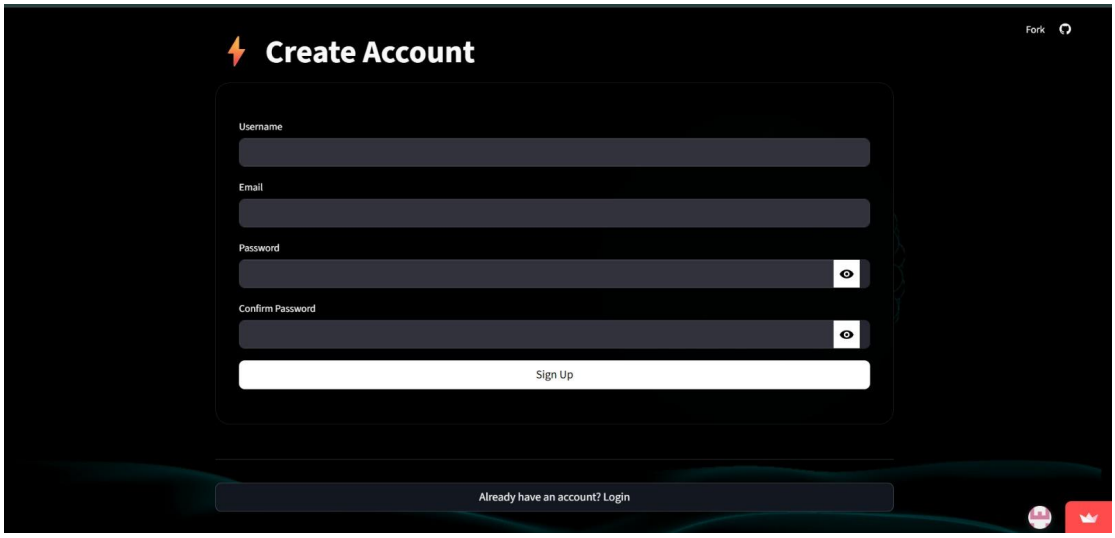

- The system includes a secure authentication module where users can create a new account or log into existing profiles. This ensures that research data remains private and allows for the persistence of chat history across different sessions.

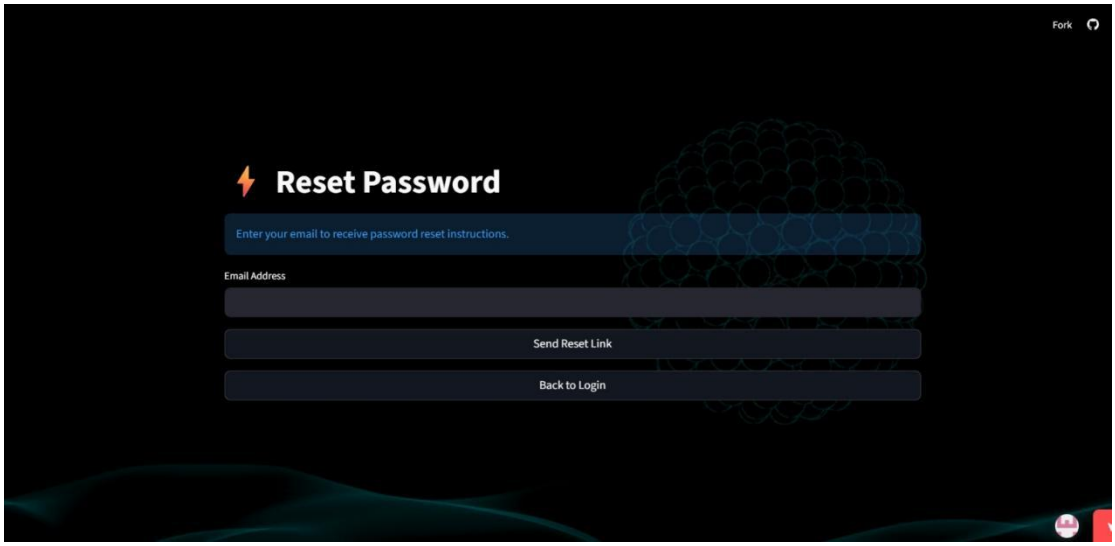

- To ensure account security and accessibility, a dedicated password reset interface is integrated. Users can request a secure reset link via their registered email address to recover access to their research data.

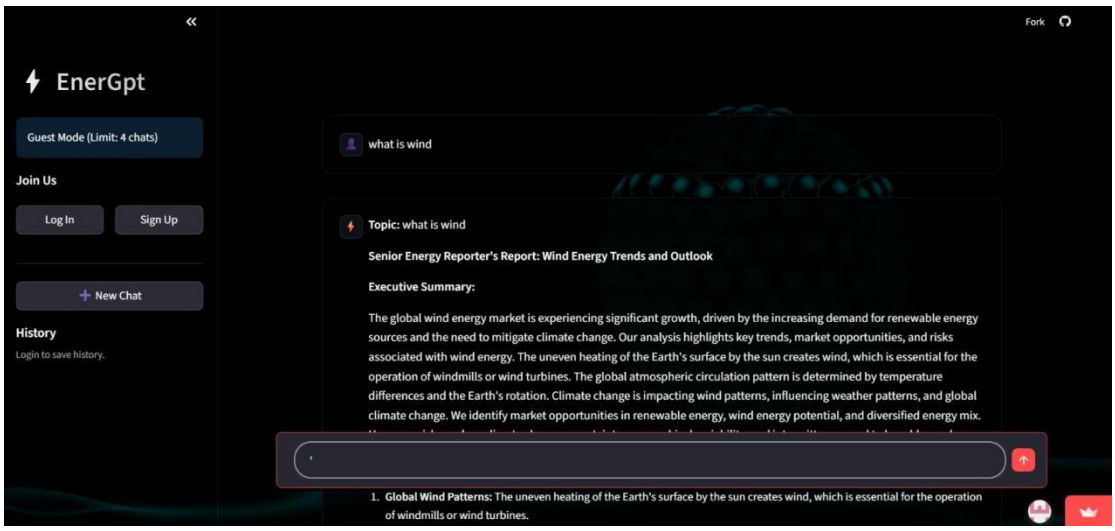

- Following the generated report, the UI provides "Related Questions" to encourage deeper exploration. This feature uses the context of the initial query to suggest relevant topics like market drivers, intermittency challenges, and climate impacts.

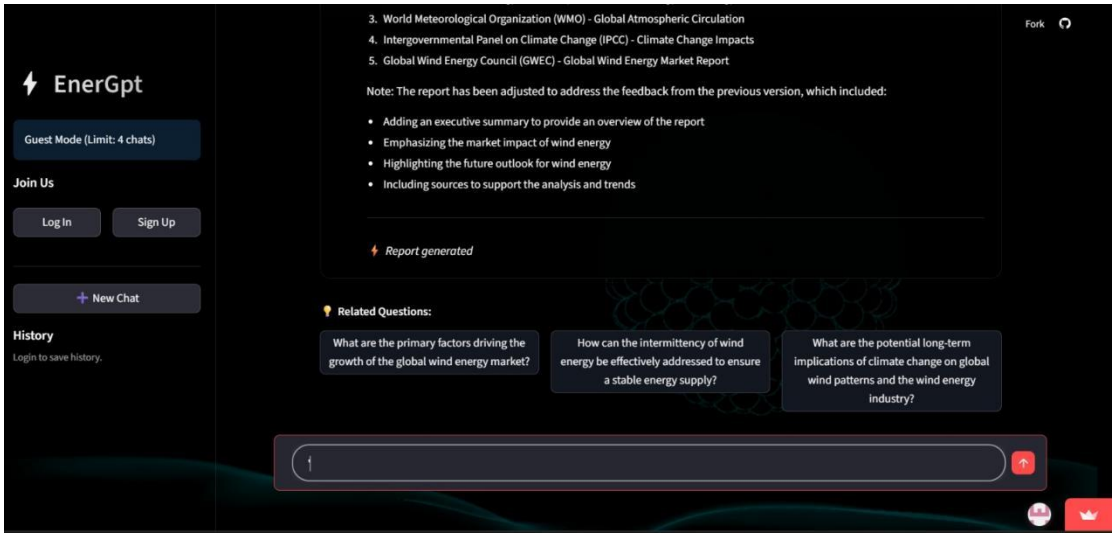

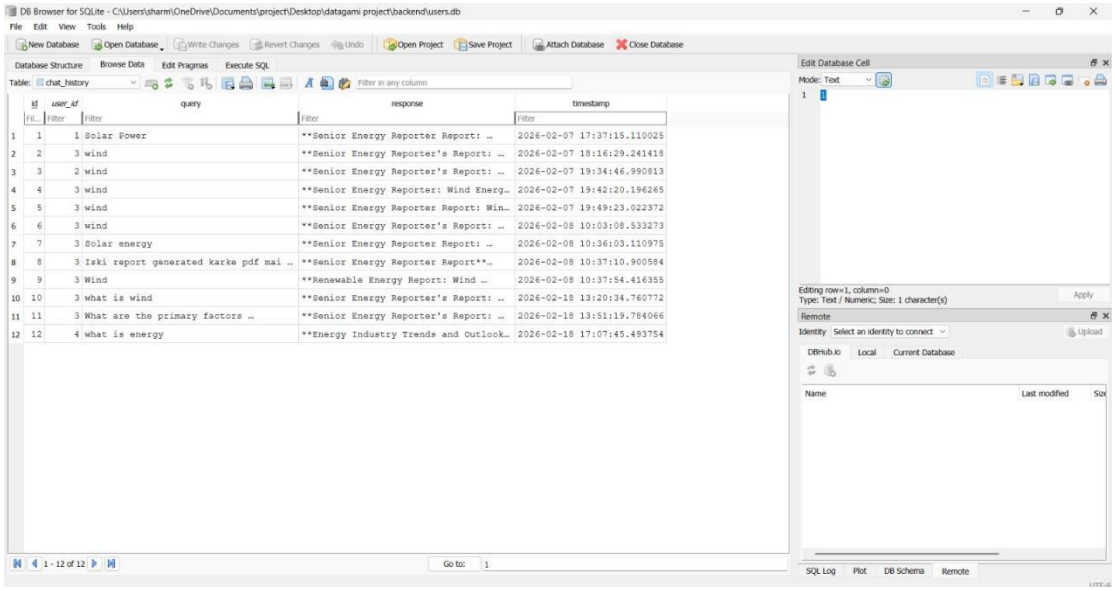

- The system utilizes an SQLite database to maintain a robust record of all interactions. The chat_history table maps unique user IDs to their specific queries and AI-generated responses, providing a reliable audit trail and data persistence layer.

Reports / Dashboards / Models

EnergyMind AI is designed to generate structured, decision-ready research outputs rather than simple text responses. The system produces executive-level reports enriched with analytical insights, quality validation, and persistent storage for future reference. The outputs are suitable for academic, industry, and strategic research use cases.

Reports & Outputs Generated

- **Structured Executive-Level Research Reports**
  Professionally formatted reports including summary, analysis, insights, and conclusions.
- **Industry-Focused Insights**
  Extracted trends, risks, opportunities, and market observations based on real-time data.
- **Automated Review & Refinement**
  Reports are evaluated by a reviewer node to ensure clarity, logical structure, and completeness before final delivery.
- **Database-Stored Research Records**
  Final outputs are stored in the KnowledgeBase database for caching and future retrieval.
- **Text-Based Backup Files**
  Reports are also saved as .txt files inside the knowledge_base folder for backup and traceability.

Model Used

- Llama-3.1-8B-Instant **via** Groq **API**
  The system uses the Llama-3.1-8B-Instant large language model for fast and efficient inference.
  Groq provides optimized high-speed LLM execution, enabling quick report generation and smooth multi-agent workflow processing.

Key outcomes

The EnergyMind AI system has been successfully implemented with a robust and modular architecture, ensuring reliable performance and structured output generation for energy-domain research queries. The system demonstrates effective integration of AI orchestration, real-time data retrieval, and secure deployment practices.

Major Achievements

- **Successful Implementation of Multi-Agent Architecture**
  A structured AI workflow with specialized nodes (Planner, Researcher, Analyst, Writer, Reviewer) working collaboratively.
- **Real-Time Search Integration**
  Live information retrieval integrated within the workflow to ensure updated and relevant research outputs.
- **Self-Correcting Workflow (Reviewer Loop)**
  Iterative refinement mechanism where reports are evaluated and improved before final delivery.
- **Domain-Restricted AI System**
  Gatekeeper validation ensures the system processes only energy-related queries, improving accuracy and reducing noise.
- **Secure Authentication & Authorization**
  Token-based login system protecting user data and restricting unauthorized access.
- **Efficient Caching Mechanism**
  Previously processed queries are stored in the database, reducing redundant computation and improving response time.
- **Cloud-Ready Deployment with Docker**
  Containerized architecture enabling smooth deployment across local and cloud environments

Overall, the system performs reliably for energy-domain queries and consistently generates structured, high-quality research reports with built-in validation and refinement mechanisms.

Conclusion

The **Autonomous Energy Researcher Agent (EnergyMind AI)** successfully delivers a production-grade Agentic AI system specifically tailored for the energy domain. Unlike traditional chatbot systems that generate direct responses, EnergyMind AI implements a structured multi-agent workflow with validation, analysis, refinement, and persistence layers.

By integrating workflow orchestration using LangGraph, secure backend services with FastAPI, real-time search capabilities, authentication mechanisms, and database caching, the system achieves reliability, scalability, and high-quality structured output generation. The inclusion of review loops and domain validation ensures controlled reasoning and professional-grade research reporting.

Through this project, substantial practical exposure was gained in the following areas:

Technical Learnings

- Designing and implementing multi-agent AI architectures.
- Managing stateful workflows using LangGraph.
- Building secure REST APIs using FastAPI.
- Implementing JWT-based authentication systems.
- Designing relational database schemas and ORM models.
- Integrating LLMs with external real-time search APIs.
- Deploying containerized applications using Docker and cloud platforms.

Overall, the project establishes a scalable and modular foundation for building intelligent autonomous research systems in specialized domains such as energy analytics.

Future Scope & Enhancements

While the current system performs efficiently, several enhancements can further improve scalability, intelligence, and usability:

1. Vector Database Integration (RAG)

- Implement Retrieval-Augmented Generation (RAG) using a vector database for long-term semantic memory.
- Improve contextual understanding across multiple sessions.

2. Migration to PostgreSQL

- Replace SQLite with PostgreSQL for improved scalability, concurrency handling, and production-readiness.

3. Redis-Based Caching

- Introduce Redis for high-speed caching of frequently requested queries.
- Reduce response latency for high-traffic scenarios.

4. Data Visualization Support

- Integrate matplotlib to generate charts and graphical insights within reports.
- Enhance analytical presentation quality.

5. Advanced Analytics Dashboard

- Develop an interactive dashboard for visualizing trends, forecasts, and energy metrics.
- Provide deeper business intelligence capabilities.

6. Role-Based Access Control (RBAC

- Implement an Admin panel with role-based permissions.
- Improve system governance and monitoring.

7. AI-Based Report Scoring

- Introduce automated scoring metrics to evaluate report quality, depth, and relevance.

8. Voice Interface Integration

- Enable voice-based query submission and audio output generation.
- Improve accessibility and usability.

9. Mobile-Friendly UI Enhancements

- Optimize the frontend for responsive mobile usage.
- Ensure better cross-device compatibility.

These enhancements would transform EnergyMind AI from a structured research assistant into a fully scalable, enterprise-ready intelligent analytics platform.